SOLAR-ASSISTED COMPUTE FOR SMALL AI LABS

Intelligence powered by sunlight.

An accelerator workstation, a solar array and storage, sized and scheduled together, so small labs can fine-tune and run their own models on their own hardware.



RAISING

$5M Pre-Seed

FOUNDER

Brian Simiyu
bw@circleplus.xyz

WEB

combinelabs.tech

COMBINE, A MADE OF INTELLIGENCE COMPANY

NOVEMBER 2026

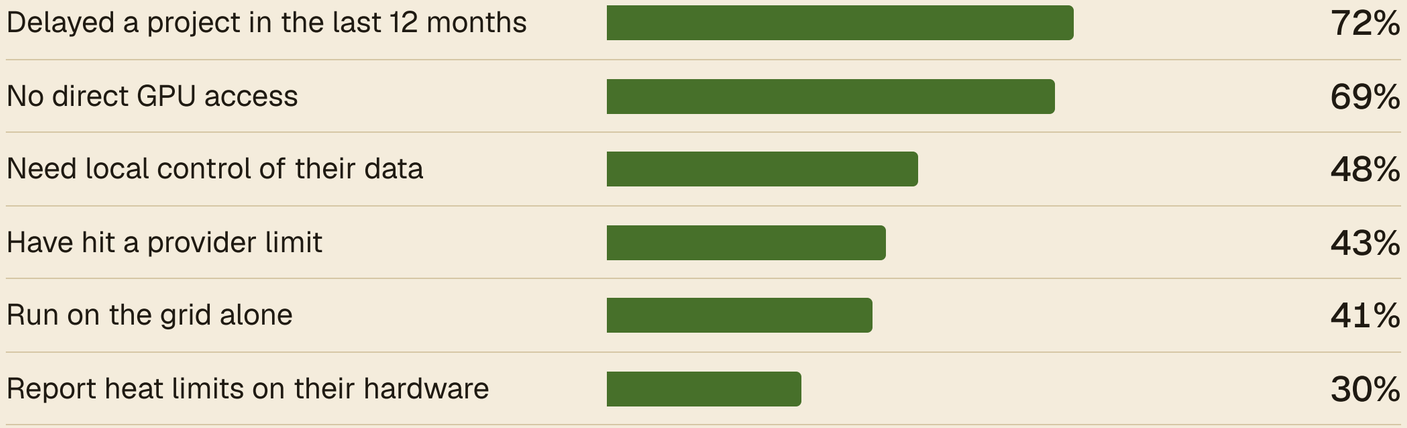
Small AI labs are held back by access, not talent.

Two in three AI builders have no GPU of their own. The cost shows up as delayed and abandoned work.

EXHIBIT 1

Delay and access top the list of what blocks small labs

Share of respondents, % (n = 600)



14.7 h

Median wait for a compute job

Queues and quotas hold work back.

3 h

Lost to power outages each week

The median respondent has grid power 22 hours a day.

SOURCE: COMBINE SURVEY OF 600 AI PRACTITIONERS IN 42 COUNTRIES, FEBRUARY TO OCTOBER 2026

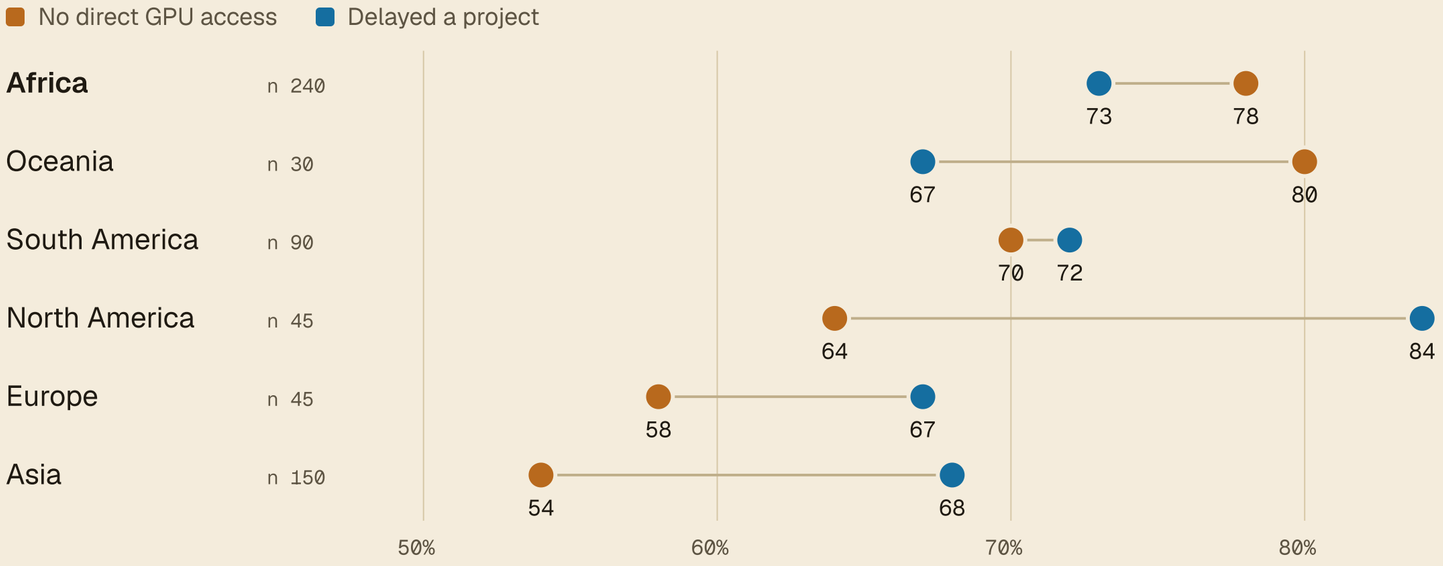
The gap is widest in Africa, where we start.

Our first customer is an independent builder or small applied research team adapting open-weight models, with a defined task and a sunny site.

EXHIBIT 2

Africa pairs the most respondents with the least GPU access

Share of respondents by region, %



240

Respondents in Africa, the largest group

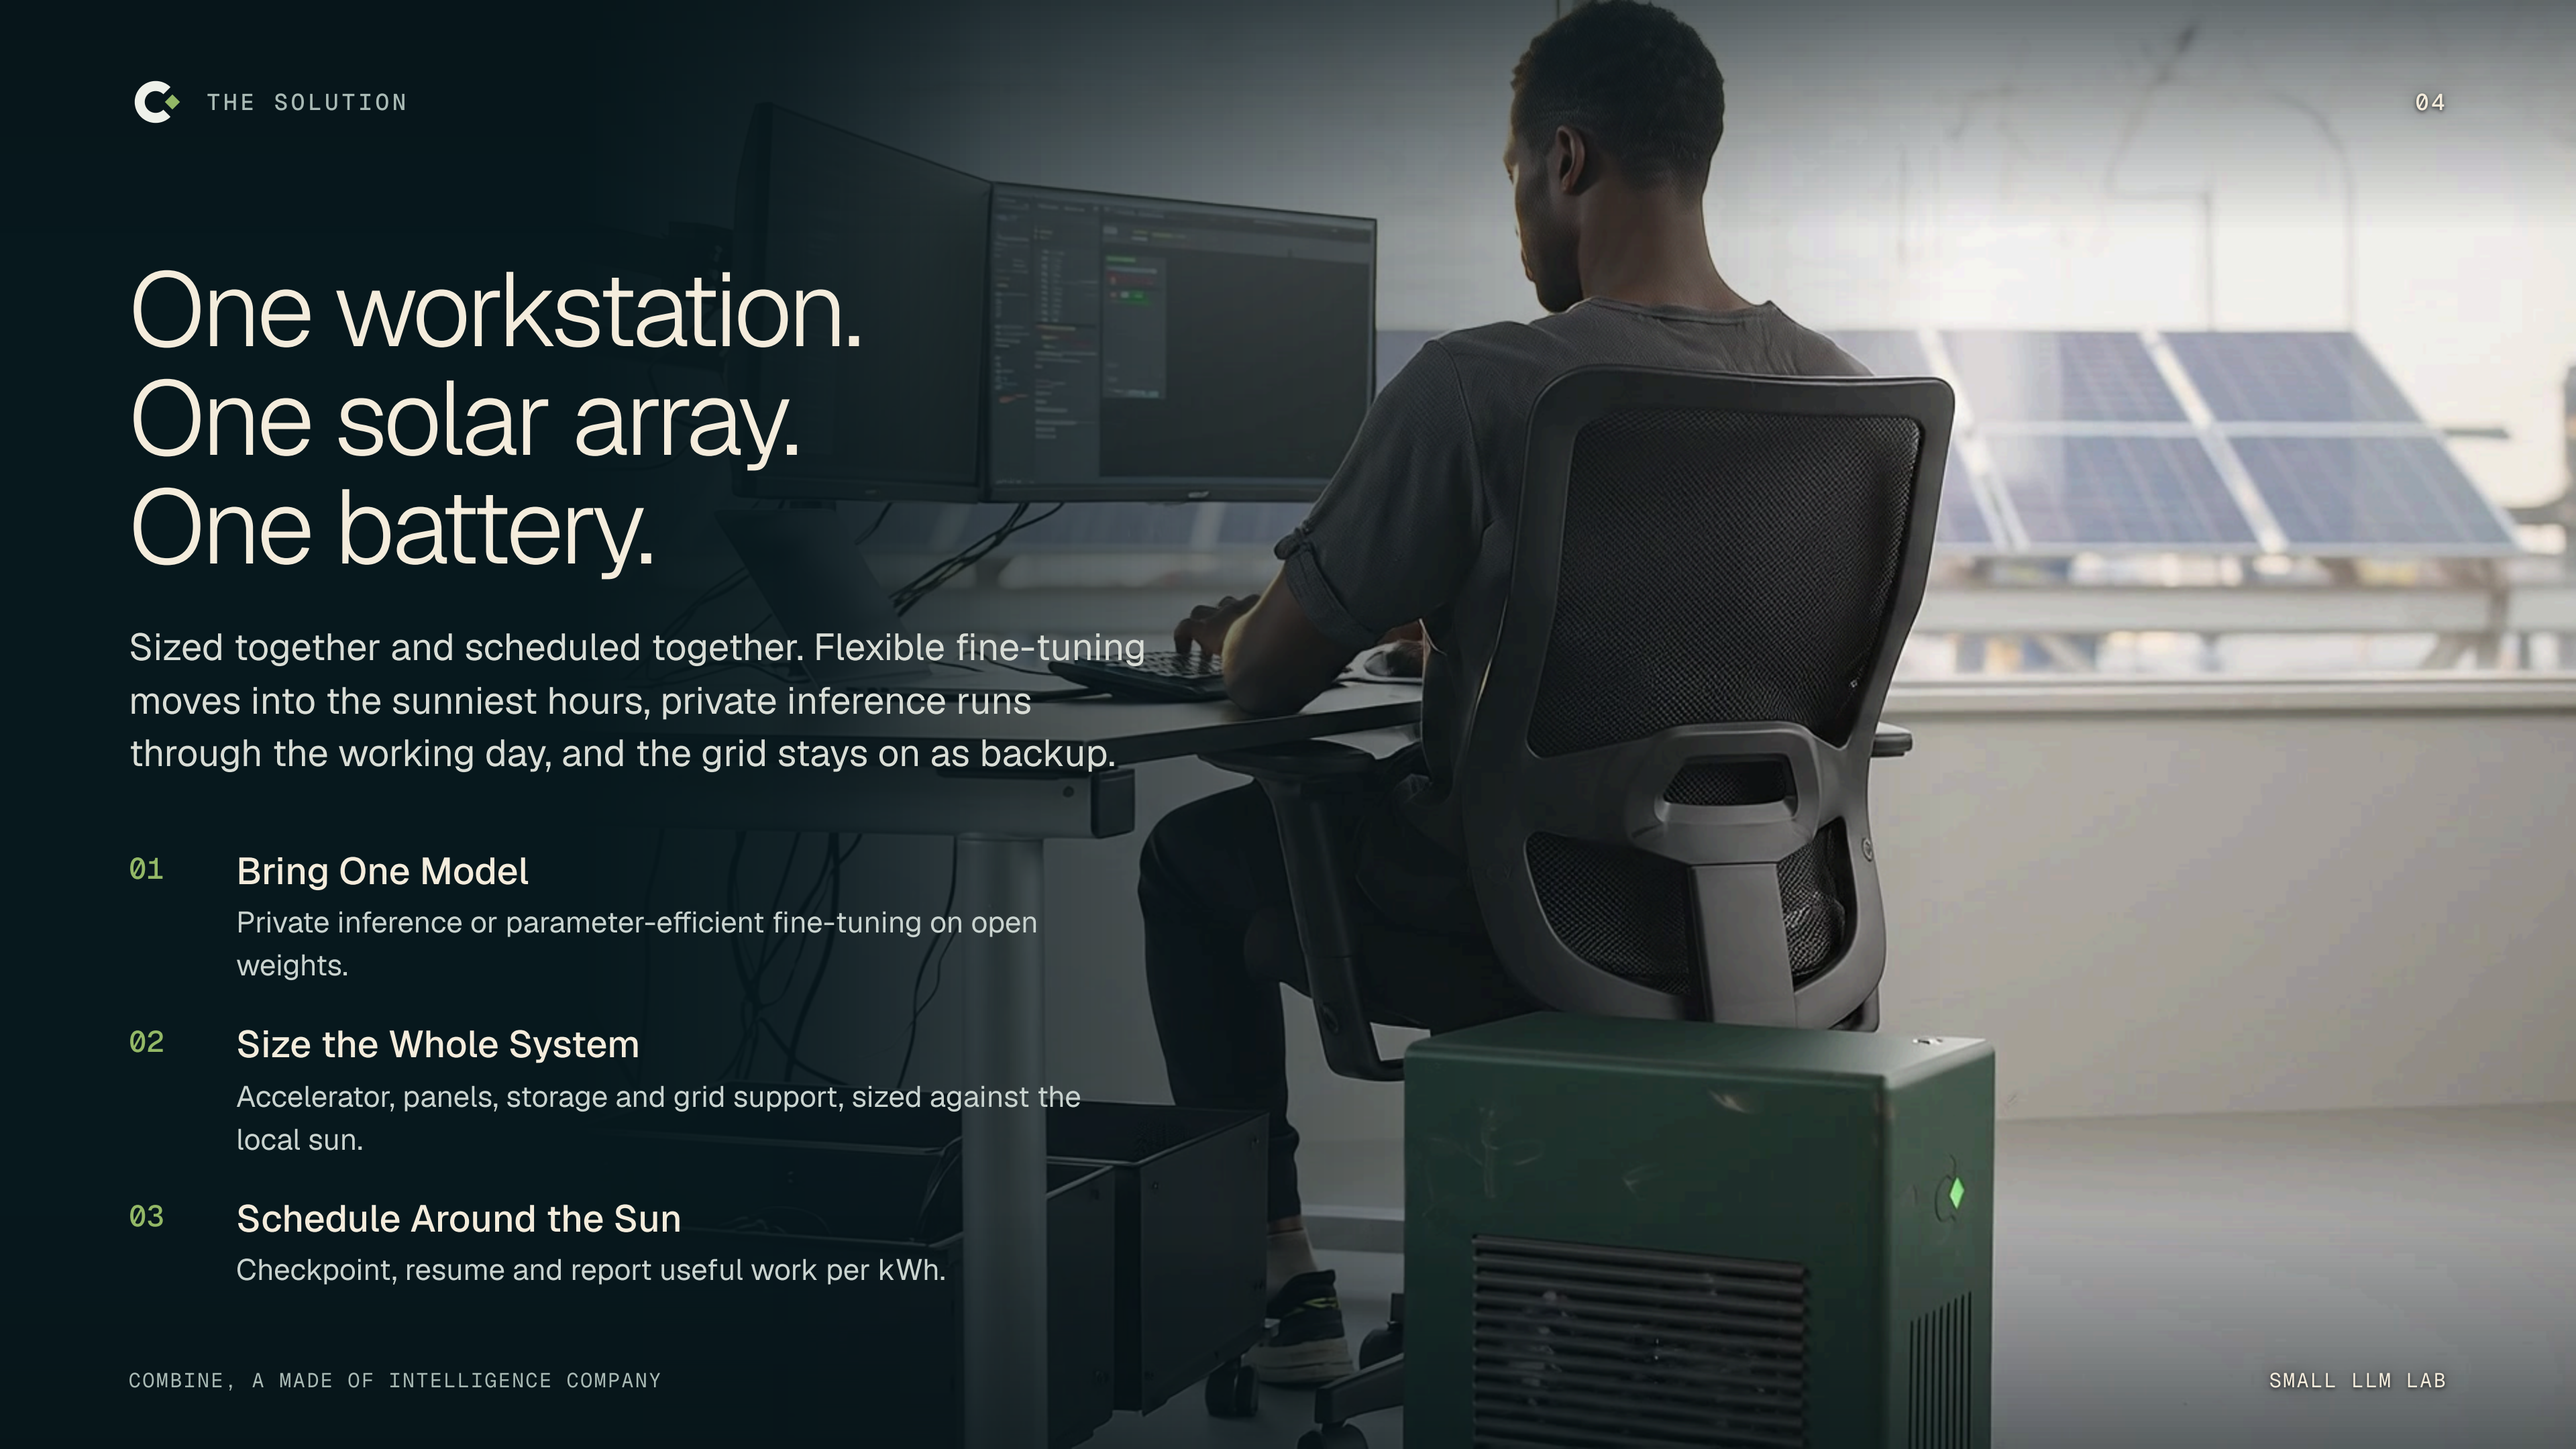
78% have no direct GPU and wait a median 15.8 hours for a job.

39%

Do LoRA or QLoRA fine-tuning

Their work is adaptation and local inference on Llama, Mistral, Gemma and Qwen, a fit for one workstation.

SOURCE: COMBINE SURVEY OF 600 AI PRACTITIONERS IN 42 COUNTRIES, 1 FEBRUARY TO 1 OCTOBER 2026. OCEANIA HAS 30 RESPONDENTS



One workstation. One solar array. One battery.

Sized together and scheduled together. Flexible fine-tuning moves into the sunniest hours, private inference runs through the working day, and the grid stays on as backup.

01 **Bring One Model**

Private inference or parameter-efficient fine-tuning on open weights.

02 **Size the Whole System**

Accelerator, panels, storage and grid support, sized against the local sun.

03 **Schedule Around the Sun**

Checkpoint, resume and report useful work per kWh.

The node. Three parts, one system.

About 2 kWp of panels, 5 to 15 kWh of storage and one accelerator workstation.



01 Workstation

One 600 W class accelerator for parameter-efficient fine-tuning and private local inference.



02 Storage & Inverter

5 to 15 kWh of usable lithium storage, with grid support for shortfalls.



03 Outdoor PV

Four 500 W panels on a fixed rack, sized to the local sun.

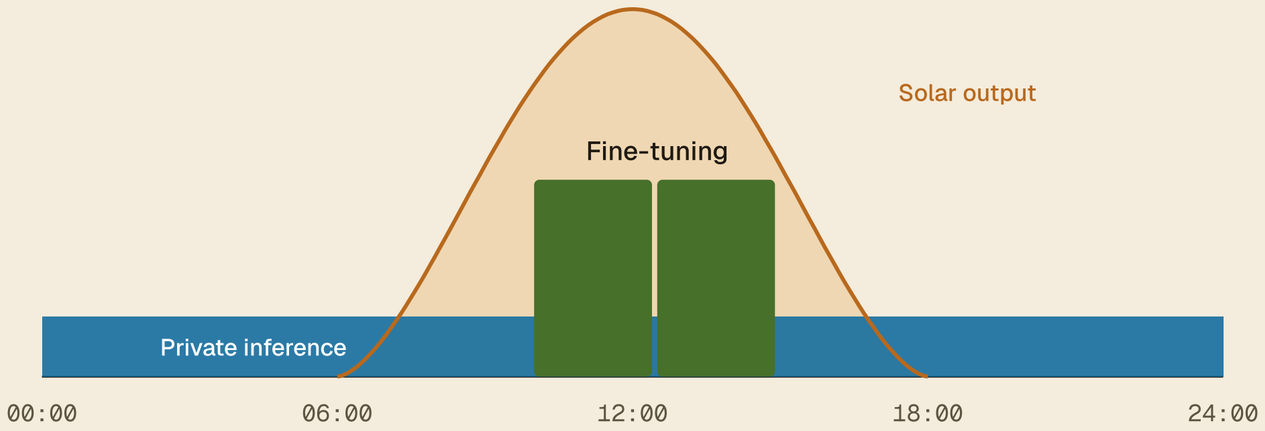
The heaviest work runs in the brightest hours.

Fine-tuning queues for the solar peak, inference runs through the day, and every run checkpoints so a drop in power never loses work.

EXHIBIT 3

A day on one node

Solar output and scheduled work across 24 hours



~97%

Of a year's load from the sun

MODELLED

A 2 kWp array and 5 kWh of storage on a typical sunny site.

62%

Less grid import, best case

MODELLED

From sun-aware scheduling alone.

Software that ships with every node

Energy-aware scheduling, checkpoint and resume, device health, and a report of useful work per kWh.

ENERGY MODEL CHECKED AGAINST PVLIB AND NREL SAM, NOVEMBER 2026

One idea, three scales.

The lab is the first product. The same sizing and scheduling carries to a logistics yard and a data center campus.



01 Small LLM Lab

One workstation, one array, one battery. Fine-tune and run private models on hardware the lab controls.



02 Bay Yard

A pole-mounted camera node with its own panel and battery reads bay status on a small vision model.



03 Data Center

Accelerator racks beside a solar field, storage and a utility connection, with firm power carrying nights.

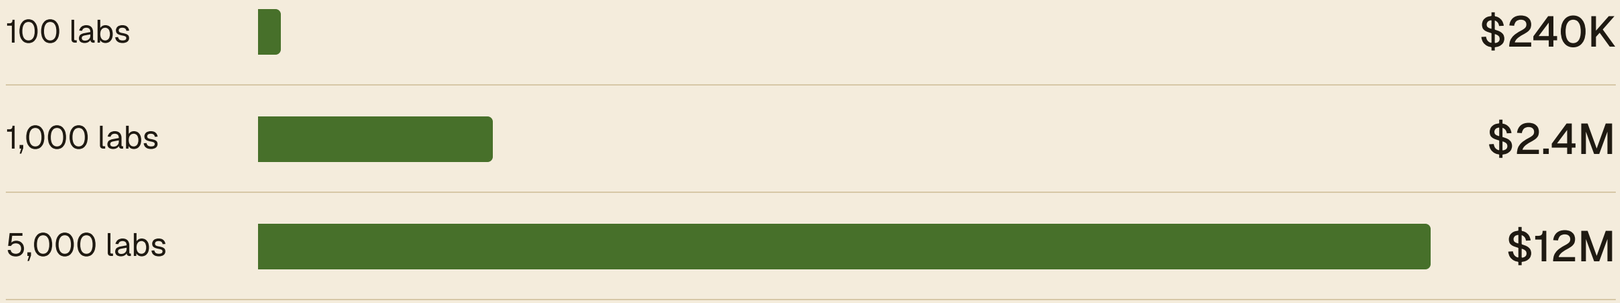
Start with small labs. Grow into the grid's edge.

Our beachhead is small labs that need local control and have sun to spare, starting in Africa. The same node and software carry to yards and campuses.

EXHIBIT 4

Software revenue scales with the number of labs

Annual recurring software revenue at $200 per lab per month, US$



Deployment and integration revenue comes on top.

Paid evaluation first. Recurring software after.

We sell to the lab lead or technical founder who owns the workload, the budget and the installation decision.

DEPLOYMENT

$1,000 to $3,000

Integration fee per lab for a scoped evaluation or installation. Hardware and electrical work quoted separately.

RECURRING SOFTWARE

$200 per lab per month

Energy-aware scheduling, device health and reporting of useful work per kWh.

01 Discovery

Qualify the model, site and budget, starting with the 131 labs that offered to host a pilot.

02 Paid Evaluation

A four-week test against the lab's current setup, with whole-workstation power measured.

03 Repeat Deployment

Turn each pilot into a standard scope, a subscription and a referral. Sales cycle of 1 to 3 months.

The benchmark is the lab's best alternative.

Combine brings compute, power and scheduling together on the customer's own site.

BUYER DECISION	Cloud GPU Rental	Self-Built Workstation	System Integrator	
Location and control	Provider infrastructure	Customer site	Customer site	Customer site
Initial investment	Rent capacity	Buy hardware	Hardware and service	Hardware and integration
Power	Provider operates it	Customer designs it	Depends on scope	Solar, storage and grid, sized as one
Workload scheduling	Provider and user tools	User-managed	Depends on scope	Energy-aware, around the sun
Main trade-off	Recurring cost, data leaves the site	Engineering effort	Service dependence	Needs a site with sun

Defensibility: repeatable integrations and operating data from every node.

CLOUD EXAMPLE: AWS EC2 G5. CONFIGURATIONS VARY BY PROVIDER

A pilot pipeline before the first node ships.

131

Labs in the pilot pipeline

Teams that offered a site and a workload to host a pilot.

NVIDIA

Inception member

Synopsys

Chip partner, pending approval

Helping us build our energy management chip.

ROADMAP



DONE
Product Design
Node, hardware and software designed for one lab.

DONE
Sizing Engine
Energy model and Solar Atlas live at combinelabs.tech.

DECEMBER 2026
First Node
Built and power-measured on real fine-tuning runs.

MONTH 6
3 Pilots
Supervised installs with pipeline labs.

MONTH 12
10 Paying Labs
Subscriptions live, standard install scope.

MONTH 24
Repeat Installs
Renewals, referrals and the first yard sites.

Built on NVIDIA. Headed for silicon.

NVIDIA Inception

Local inference and fine-tuning on an NVIDIA accelerator matched to model size.

Synopsys

Helping us build our energy management chip. The partnership is pending approval.

Inside the Chip

Our energy management and scheduling block, in simulation today. Silicon follows measured workloads.



CHIP IMAGE: CONCEPT RENDER OF A PROPOSED DESIGN

Founder-led, with hardware and energy depth.



Brian Simiyu

FOUNDER & CEO

Third-time technical founder and Techstars SF alumnus. His previous startup reached over $100K MRR.



Natalie Mutunga

CHIEF HARDWARE ENGINEER

MSc in Energy & Economic Finance, University of Aberdeen.

techstars_

NVIDIA INCEPTION PROGRAM

1495 UNIVERSITY OF ABERDEEN

Raising $5M pre-seed.

Seeking investor introductions now.

24 months

Runway

$4.5M operating at about $187,500 a month, plus $500K contingency.

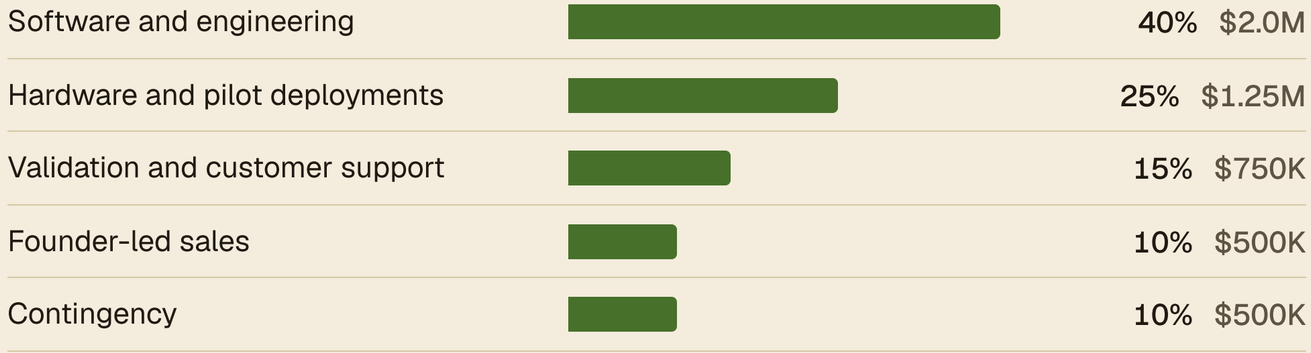
FUNDED MILESTONES

Month 6	Baseline measured, 3 supervised pilots
Month 12	10 paying labs
Month 24	Repeat installs and renewal data

EXHIBIT 5

Use of funds

Share of the $5M round, % and US$



Intelligence powered by sunlight.

Solar-assisted compute for small AI labs. Raising $5M pre-seed to put the first nodes in the field.



BRIAN SIMIYU
bw@circleplus.xyz

LINKEDIN
linkedin.com/in/bw-simiyu

WEB
combinelabs.tech